


DEPLOYABLE INFERENCE



OFFLINE TRAINING AND EVALUATION

Excluded from model input: true β · true y/n · configuration ID · seed · repeat_id · candidate δ

Training labels (formal MC only)

raw per-sample vector $[\ell_i(\delta_0), \dots, \ell_i(\delta_{25})]$

5-fold full-combination holdout

Evaluation only — not a deployment input

run MDM at $\hat{\delta}_i \rightarrow$ true selected-loss $\ell_i(\hat{\delta}_i)$

aggregate after selection: $J_1 = \sqrt{\text{mean}_i \ell_i(\hat{\delta}_i)}$